

Sec_1_Homework_9

March 6, 2024

1 0.) Import and Clean data

```
[81]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
```

```
[49]: #drive.mount('/content/gdrive/', force_remount = True)
df = pd.read_csv("Country-data.csv", sep = ",")
df.set_index('country', inplace = True)
df
```

```
[49]:
```

	child_mort	exports	health	imports	income	inflation	\
country							
Afghanistan	90.2	10.0	7.58	44.9	1610	9.44	
Albania	16.6	28.0	6.55	48.6	9930	4.49	
Algeria	27.3	38.4	4.17	31.4	12900	16.10	
Angola	119.0	62.3	2.85	42.9	5900	22.40	
Antigua and Barbuda	10.3	45.5	6.03	58.9	19100	1.44	
...		
Vanuatu	29.2	46.6	5.25	52.7	2950	2.62	
Venezuela	17.1	28.5	4.91	17.6	16500	45.90	
Vietnam	23.3	72.0	6.84	80.2	4490	12.10	
Yemen	56.3	30.0	5.18	34.4	4480	23.60	
Zambia	83.1	37.0	5.89	30.9	3280	14.00	

	life_expec	total_fer	gdpp
country			
Afghanistan	56.2	5.82	553
Albania	76.3	1.65	4090
Algeria	76.5	2.89	4460
Angola	60.1	6.16	3530
Antigua and Barbuda	76.8	2.13	12200
...
Vanuatu	63.0	3.50	2970
Venezuela	75.4	2.47	13500

Vietnam	73.1	1.95	1310
Yemen	67.5	4.67	1310
Zambia	52.0	5.40	1460

[167 rows x 9 columns]

```
[91]: X = df
      scaler = StandardScaler()
      X_scaled = scaler.fit_transform(X)
```

2 1.) Fit a kmeans Model with any Number of Clusters

```
[93]: kmeans = KMeans(n_clusters=5, n_init=1).fit(X_scaled)
      kmeans
```

```
[93]: KMeans(n_clusters=5, n_init=1)
```

3 2.) Pick two features to visualize across

```
[59]: X.columns
```

```
[59]: Index(['child_mort', 'exports', 'health', 'imports', 'income', 'inflation',
           'life_expec', 'total_fer', 'gdpp'],
          dtype='object')
```

```
[90]: import matplotlib.pyplot as plt

      x1_index = 0
      x2_index = 4

      # Apply KMeans clustering
      kmeans = KMeans(n_clusters=3) # Example: assuming 3 clusters
      kmeans.fit(X_scaled)

      # Obtain cluster labels
      #cluster_labels = kmeans.labels_

      # Plot the scatter plot of data points with KMeans clustering labels
      scatter = plt.scatter(X_scaled[:, x1_index], X_scaled[:, x2_index], cmap = '
      ↪ viridis', c=kmeans.labels_, label='Clusters')

      # Plot the cluster centers
      centers = plt.scatter(kmeans.cluster_centers_[:, x1_index], kmeans.
      ↪ cluster_centers_[:, x2_index], marker='o', color='black', s=100,
      ↪ label='Centers')
```

```

# Set labels and title
plt.xlabel(X.columns[x1_index])
plt.ylabel(X.columns[x2_index])
plt.title('Scatter Plot of Customers')

# Generate legend
plt.legend()


# Display grid
plt.grid()

# Show the plot
plt.show()

```

/Users/trujillo/anaconda3/lib/python3.11/site-packages/sklearn/cluster/_kmeans.py:1416: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning

```
super()._check_params_vs_input(X, default_n_init=10)
```




- 4 3.) Check a range of k-clusters and visualize to find the elbow.
Test 30 different random starting places for the centroid means

```
[87]: WCSSs = []  
Ks = range(2, 15)  
for k in Ks:  
    kmeans = KMeans(n_clusters=k, n_init=30).fit(X_scaled)  
    WCSSs.append(kmeans.inertia_)
```

```
[ ]:
```

```
[88]: plt.plot(Ks, WCSSs)  
plt.xlabel('Number of Clusters (k)')  
plt.ylabel('Within-Cluster Sum of Squares (WCSS)')  
plt.title('Elbow Method')
```


```
[88]: Text(0.5, 1.0, 'Elbow Method')
```



- 5 4.) Use the above work and economic critical thinking to choose a number of clusters. Explain why you chose the number of clusters and fit a model accordingly.

```
[89]: plt.plot(Ks,WCSSs)
plt.xlabel('# of clusters')
plt.ylabel('WCSS')
```

```
[89]: Text(0, 0.5, 'WCSS')
```



Even if the optimal number of clusters cannot be seen in this plot (there is no kink), economic rationale would tell us to ask for two clusters. That way we only divide the sample into advanced and developed economies. There might be many more options, but to keep it simple we will select only 2.


- 6 6.) Do the same for a silhouette plot

```
[94]: from sklearn.metrics import silhouette_score
```

```
[95]: SSs = []
      Ks = range(2, 15)
      for k in Ks:
          kmeans = KMeans(n_clusters=k, n_init=30).fit(X_scaled)
          sil = silhouette_score(X_scaled, kmeans.labels_)
          SSs.append(sil)
```

```
[41]: plt.plot(Ks,SSs)
      plt.xlabel('Number of Clusters (k)')
      plt.ylabel('Silhouette Score')
      plt.title('Silhouette Method')
```

```
[41]: Text(0.5, 1.0, 'Silhouette Method')
```



7 7.) Create a list of the countries that are in each cluster. Write interesting things you notice.

```
[96]: kmeans = KMeans(n_clusters=2, n_init=30).fit(X_scaled)
```

```
[98]: preds = kmeans.labels_  
  
df['preds'] = preds
```

```
[99]: print('Cluster 1: ', df.index[preds==0] )
```

```
Cluster 1: Index(['Albania', 'Algeria', 'Antigua and Barbuda', 'Argentina',  
'Armenia',  
                'Australia', 'Austria', 'Azerbaijan', 'Bahamas', 'Bahrain', 'Barbados',  
                'Belarus', 'Belgium', 'Belize', 'Bhutan', 'Bosnia and Herzegovina',  
                'Brazil', 'Brunei', 'Bulgaria', 'Canada', 'Cape Verde', 'Chile',  
                'China', 'Colombia', 'Costa Rica', 'Croatia', 'Cyprus',  
                'Czech Republic', 'Denmark', 'Dominican Republic', 'Ecuador',  
                'El Salvador', 'Estonia', 'Fiji', 'Finland', 'France', 'Georgia',  
                'Germany', 'Greece', 'Grenada', 'Hungary', 'Iceland', 'Iran', 'Ireland',  
                'Israel', 'Italy', 'Jamaica', 'Japan', 'Jordan', 'Kazakhstan', 'Kuwait',  
                'Latvia', 'Lebanon', 'Libya', 'Lithuania', 'Luxembourg',  
                'Macedonia, FYR', 'Malaysia', 'Maldives', 'Malta', 'Mauritius',  
                'Moldova', 'Montenegro', 'Morocco', 'Netherlands', 'New Zealand',  
                'Norway', 'Oman', 'Panama', 'Paraguay', 'Peru', 'Poland', 'Portugal',  
                'Qatar', 'Romania', 'Russia', 'Saudi Arabia', 'Serbia', 'Seychelles',  
                'Singapore', 'Slovak Republic', 'Slovenia', 'South Korea', 'Spain',  
                'Sri Lanka', 'St. Vincent and the Grenadines', 'Suriname', 'Sweden',  
                'Switzerland', 'Thailand', 'Tunisia', 'Turkey', 'Ukraine',  
                'United Arab Emirates', 'United Kingdom', 'United States', 'Uruguay',  
                'Venezuela', 'Vietnam'],  
              dtype='object', name='country')
```

```
[100]: print('Cluster 2: ', df.index[preds==1] )
```

```
Cluster 2: Index(['Afghanistan', 'Angola', 'Bangladesh', 'Benin', 'Bolivia',  
'Botswana',  
                'Burkina Faso', 'Burundi', 'Cambodia', 'Cameroon',  
                'Central African Republic', 'Chad', 'Comoros', 'Congo, Dem. Rep.',  
                'Congo, Rep.', 'Cote d'Ivoire', 'Egypt', 'Equatorial Guinea', 'Eritrea',  
                'Gabon', 'Gambia', 'Ghana', 'Guatemala', 'Guinea', 'Guinea-Bissau',  
                'Guyana', 'Haiti', 'India', 'Indonesia', 'Iraq', 'Kenya', 'Kiribati',  
                'Kyrgyz Republic', 'Lao', 'Lesotho', 'Liberia', 'Madagascar', 'Malawi',  
                'Mali', 'Mauritania', 'Micronesia, Fed. Sts.', 'Mongolia', 'Mozambique',  
                'Myanmar', 'Namibia', 'Nepal', 'Niger', 'Nigeria', 'Pakistan',  
                'Philippines', 'Rwanda', 'Samoa', 'Senegal', 'Sierra Leone',  
                'Solomon Islands', 'South Africa', 'Sudan', 'Tajikistan', 'Tanzania',  
                'Timor-Leste', 'Togo', 'Tonga', 'Turkmenistan', 'Uganda', 'Uzbekistan',
```

```
'Vanuatu', 'Yemen', 'Zambia'],
dtype='object', name='country')
```

```
[ ]: ##### Write an observation
```

It seems that the clustering was roughly fine. For example, in this case it is clear that cluster 1 refers to the advanced economies because we have countries as United Kingdom, United States, Netherlands, Germany, among others. However, maybe the threshold was a little bit fuzzy because we also get El Salvador, Venezuela, and Vietnam in this cluster. For cluster 2 (developing countries) all selections seem right by inspection.

#8.) Create a table of Descriptive Statistics. Rows being the Cluster number and columns being all the features. Values being the mean of the centroid. Use the nonscaled X values for interpretation

```
[77]: df.groupby('preds').mean()
```

```
[77]:
```

	child_mort	exports	health	imports	income	inflation	\
preds							
0	12.161616	48.603030	7.314040	49.121212	26017.171717	5.503545	
1	76.280882	30.198515	6.090147	43.642146	4227.397059	11.098750	

	life_expec	total_fer	gdpp
preds			
0	76.493939	1.941111	20507.979798
1	61.910294	4.413824	1981.235294

```
[78]: df.groupby('preds').std()
```

```
[78]:
```

	child_mort	exports	health	imports	income	inflation	\
preds							
0	8.523122	30.116032	2.716652	26.928785	20441.749847	6.957187	
1	38.076068	18.201742	2.645319	19.323451	4890.581414	13.682630	

	life_expec	total_fer	gdpp
preds			
0	3.735757	0.486744	20578.727127
1	6.897418	1.285590	2528.509189

8 9.) Write an observation about the descriptive statistics.

As thought before, the first cluster refers to the advanced economies since we can see lower child mortality, higher exports, much higher income and life expectancy. Also, its standard error is much lower across metrics (except exports and imports) so we can argue that the developing economies are still very far away from each other, while the advanced might be reaching some stationary state.

```
[ ]:
```